

What is layerwise relevance propagation

Layerwise relevance propagation (LRP) is a propagation-based explanation method that calculate the **quantitative contribution** of a given neuron to the prediction.

The **global LRP plot** identifies the important features and their contribution **in all cases**.

The **local LRP plot** displays the impact of each feature on the model outcome **for an individual example**.




Fig. 1 An example of the global LRP plot. The length of the bar chart represents the contribution of different features.




Fig. 2 An example of the local LRP plot. The length of the bar chart represents the contribution of different features. Red represents positive contribution and blue represents negative contribution.